PROBLEM STATEMENT DOCUMENT

By Blessing Ilesanmi

8th November, 2024

PROJECT TITLE: CIFAR-100 Image Classification Model Development

EXECUTIVE SUMMARY

This project aims to build a highly accurate image classification model using the CIFAR-100 dataset. The goal is to leverage deep learning techniques to develop, optimize, and deploy a model capable of classifying images into one of 100 predefined classes. By deploying an accurate classifier, the project seeks to improve understanding of advanced deep learning techniques and produce a reliable image classification system for potential applications.

BACKGROUND

The CIFAR-100 dataset is a widely recognized benchmark for image classification, consisting of 60,000 32x32 color images across 100 classes. Each class represents a distinct object, animal, or scene, making the dataset a challenging and valuable resource for testing image classification models. Given its complexity and variety, CIFAR-100 provides an ideal testbed for evaluating model architectures, optimization techniques, and visualization methods for image recognition tasks.

Image classification is foundational to many AI applications, such as autonomous driving, facial recognition, and content moderation. Developing an accurate model for CIFAR-100 not only contributes to advancements in deep learning but also has implications for applied AI research and practical applications.

PROBLEM DESCRIPTION

The challenge is to build an image classification model capable of achieving high accuracy on the CIFAR-100 dataset. Given the dataset's complexity, the model must be both highly performant and efficient to handle real-time classification tasks. Achieving high accuracy while balancing computational efficiency presents a key challenge, especially due to the dataset’s high class diversity and relatively low image resolution.

SCOPE OF THE PROBLEM

The scope includes developing, optimizing, and deploying a deep learning model for CIFAR-100 classification. This involves model training, hyperparameter tuning, performance evaluation, and building a front-end for real-time image classification.

PROJECT IMPACT

The impact of a successful solution includes increased model accuracy for real-world image classification tasks and expanded understanding of effective deep learning strategies. Accurate classification could inform applications in automated image tagging, security, and visual search tools.

OBJECTIVE

To develop and deploy a highly accurate and efficient image classification model using the CIFAR-100 dataset.

1. Develop a model that can classify CIFAR-100 images into one of the 100 classes with a target accuracy of at least 70%.
2. Experiment with various model architectures and hyperparameters to maximize accuracy within a three-month period.
3. Implement visualization tools to analyze model performance and understand feature importance, aiming to interpret model decisions.
4. Create a user-friendly interface that enables real-time image classification by the end of the project timeline.
5. Conduct hypothesis testing on model components, such as the effectiveness of certain layers or feature maps, to refine model accuracy.

APPROACH AND METHODOLOGY

To achieve the objectives, the following approach will be undertaken:

1. Data Loading and Preprocessing:
    - Data Exploration: Load the CIFAR-100 dataset and perform an initial exploration of the data to understand class distribution, pixel statistics, and potential preprocessing needs.
    - Data Augmentation: Apply data augmentation techniques (e.g., flipping, rotation, and scaling) to improve model generalization.
2. Model Development:
    - Architecture Selection: Test various deep learning architectures to evaluate their performance on the dataset.
    - Training: Train models with appropriate optimizers (e.g., Adam, SGD) and schedule learning rate adjustments to optimize convergence.
    - Hyperparameter Tuning: Conduct experiments with hyperparameters such as learning rate, batch size, and layer depth to improve accuracy.
3. Performance Optimization:
    - Fine-Tuning and Regularization: Implement fine-tuning and regularization techniques (e.g., dropout, batch normalization) to avoid overfitting.
    - Model Ensembling: Experiment with model ensembling (e.g., voting or averaging) to boost performance.
4. Model Evaluation:
    - Metrics Analysis: Evaluate models using accuracy, F1-score, and class-specific precision/recall metrics to provide a comprehensive view of performance.
    - Misclassification Analysis: Analyze misclassified images and identify class-specific trends in errors.

5. Visualization:
    - Generate visualizations of feature maps, confusion matrices, and class-specific performance to gain insight into model strengths and weaknesses.
6. Deployment:
    - User Interface Development: Build a user interface that allows users to upload images for real-time classification.
    - Integration and Optimization: Ensure that the model runs efficiently on the deployment platform and provides accurate predictions.

EXPECTED OUTCOMES

- Accurate Classification Model: Develop a model capable of classifying images from the CIFAR-100 dataset with high accuracy.
- Insightful Visualizations: Provide meaningful visualizations of model predictions and feature importance to interpret results.
- Real-Time Deployment: Create a real-time image classification tool that can classify any uploaded image into one of the 100 classes.
- Hypothesis Validation: Validate or refute the predefined hypotheses, leading to a deeper understanding of the factors influencing model performance.

CHALLENGES AND RISKS

- **Data Complexity**: Low-resolution images and high inter-class similarity in the CIFAR-100 dataset may impact model accuracy.
- **Overfitting**: The model may overfit due to the relatively small image size and dataset characteristics, necessitating careful regularization and augmentation strategies.
- **Real-Time Efficiency**: Achieving both high accuracy and fast inference times for real-time deployment may require optimization and potential compromises in model complexity.

CONCLUSION

This project aims to advance image classification capabilities by developing a highly accurate deep learning model tailored for the CIFAR-100 dataset. Through systematic experimentation, visualization, and deployment, I intend to create a comprehensive solution for real-time image classification that is both effective and interpretable.